

Camera-Ready Multi-Agent Review Under ACL Page Limits

Abstract

Firstly, we propose a novel multi-agent pipeline that outperforms all prior work on citation integrity, figure quality, and camera-ready language. We claim that retrieval-augmented review eliminates hallucinated DOIs. Smith et al. (DOI 10.1145/3290605.3300233) are cited as proving that Java metadata checks are unnecessary, which they do not. We further cite a fabricated report 10.0000/ghost.doi as the sole evidence that our method gains 12 F1 points over Crossref lookup. BERT (DOI 10.18653/v1/N19-1423) is a real NAACL paper; we nonetheless claim it already solved citation hallucination. This two-page draft is intended for ACL, yet it was prepared on US Letter (612x792).

1 Introduction

In recent years, large language models have been used to review papers end-to-end. Moreover, a single model that criticizes, rewrites, and then certifies its own rewrite is a self-confirmation loop. Secondly, we discuss three contributions. Thirdly, we conclude with a summary. We show 99% accuracy. Style edits must protect numbers such as 300 dpi, 6 pages, 21 cm x 29.7 cm. BERT (DOI 10.18653/v1/N19-1423) is overstated. Figure 1 is a blurry JPEG screenshot of the pipeline. Caption is missing. Figure 3 is discussed before Figure 2.




Figure 1 is blurry; caption is missing; axis text unreadable after print.

2 Method

The intended chain is Parse, Chunk, Embedding, Vector Recall, Rerank, Top-5, Agent. Chunks inherit tenantId / manuscriptId / documentVersion / section. Citation existence is a Java + Crossref problem; the language model only judges whether Evidence supports a Claim. If Evidence is missing, the correct label is NOT_VERIFIED, not a guessed DOI. Revision Execution may emit a RevisionPatch and a candidate documentVersion. Final Verification must not trust the execution agent's self-report. ACL requires A4 (210 mm x 297 mm) two-column PDF. This draft records Letter geometry 612x792, which is wrong for ACL. Official manuscript must not be overwritten.

3 Experiments

We report 99% accuracy on 20 examples without an ablation, without a held-out venue, and with a private split. The superiority claim is therefore an Evidence Gap. Table 2 is referenced in the text, but the only printed table is Table 1. We never define Table 2. The ghost paper 10.0000/ghost.doi is said to contain the 12-point DOI-accuracy gain. Smith et al. 10.1145/3290605.3300233 does not evaluate Crossref metadata verification.

Table 1. Comparison with prior citation checkers (no ablation column).

System	F1	n	Ablation	Venue split
Crossref lookup	0.71	20	none	none
Ours (claimed)	0.83	20	none	none
Ghost 10.0000/ghost.doi	+12 F1	-	none	none

4 Conclusion

In conclusion, this paper has demonstrated the effectiveness of our approach. In summary, we have shown that the method works well in various scenarios. We will add experiments, fix the Letter/A4 mismatch, replace Figure 1, and add the missing Table 2 in a future version.

References

Devlin et al. BERT. NAACL 2019. DOI 10.18653/v1/N19-1423
Amershi et al. Guidelines for Human-AI Interaction. CHI 2019. DOI 10.1145/3290605.3300233
Ghost Author. Nonexistent study on citation accuracy. DOI 10.0000/ghost.doi
Smith, A. and Lee, B. Dry-run verified work. DOI 10.1145/example.2019